

SDS Assignment 1 - Digital Nomad

- First of all i import the packeages that I need and download the data.
- Take a look at the data using head and info.
- Then open the files and assigning them to variables.

In []:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

In []:

```
#!wget https://sds-aau.github.io/SDS-master/M1/data/trips.csv
```

In []:

```
!wget https://sds-aau.github.io/SDS-master/M1/data/people.csv
```

In []:

```
!wget https://sds-aau.github.io/SDS-master/M1/data/countrylist.csv
```

In []:

```
tripdata = pd.read_csv('trips.csv')
peopledata = pd.read_csv('people.csv')
countrydata = pd.read_csv('countrylist.csv')
```

In []:

```
#taking a look at each file and see how the data looks
print(peopledata.head())
#tripdata.head()
countrydata.head()
#Could also use this
#peopledata.info()
#peopledata.describe()
```

```
   Unnamed: 0  ...  education_raw
0           0  ...  High School, Bachelor's Degree
1           1  ...                NaN
2           2  ...                NaN
3           3  ...                NaN
4           4  ...                NaN
```

[5 rows x 6 columns]

Out[]:

	alpha_2	region	sub_region
0	AF	Asia	Southern Asia
1	AX	Europe	Northern Europe
2	AL	Europe	Southern Europe
3	DZ	Africa	Northern Africa
4	AS	Oceania	Polynesia

Part 1 Preprocessing

Preprocessing 1A

a. Transform dates into timestamps in the Trips data.

In []:

```
# the code from 1A
#tripdata.info()
tripdata['date_end'] = pd.to_datetime(tripdata.date_end, errors = 'coerce')
tripdata['date_start'] = pd.to_datetime(tripdata.date_start, errors = 'coerce')
tripdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 46510 entries, 0 to 46509
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Unnamed: 0      46510 non-null  int64
1   username        46510 non-null  object
2   country         46473 non-null  object
3   country_code    45909 non-null  object
4   country_slug    46473 non-null  object
5   date_end        46292 non-null  datetime64[ns]
6   date_start      46499 non-null  datetime64[ns]
7   latitude        46510 non-null  int64
8   longitude       46510 non-null  int64
9   place          46510 non-null  object
10  place_slug      46473 non-null  object
dtypes: datetime64[ns](2), int64(3), object(6)
memory usage: 3.9+ MB
```

Preprocessing 1B

b. Calculate trip duration in days and create a column with a numerical value.

In []:

```
#tripdata.describe()
tripdata['duration'] = tripdata.apply(lambda row: row.date_end - row.date_start, axis=1)
tripdata.describe()
#tripdata.head()
```

Out[]:

	Unnamed: 0	latitude	longitude	duration
count	46510.000000	46510.000000	46510.000000	46282
mean	23254.500000	28.931477	11.912750	50 days 13:04:52.312346052
std	13426.424847	23.551976	78.767316	658 days 14:15:01.267546536
min	0.000000	-79.000000	-175.000000	-16831 days +00:00:00
25%	11627.250000	14.000000	-68.000000	3 days 00:00:00
50%	23254.500000	37.000000	9.000000	7 days 00:00:00
75%	34881.750000	47.000000	98.000000	23 days 00:00:00
max	46509.000000	78.000000	178.000000	74633 days 00:00:00

In []:

```
tripdata.head()
```

Out[]:

Unnamed: 0	username	country	country_code	country_slug	date_end	date_start	latitude	longitude	place	pl
------------	----------	---------	--------------	--------------	----------	------------	----------	-----------	-------	----

0	Unnamed: 0	@lewellenmichael	Mexico	country_code	country_flag	2018-06-03	2018-06-03	21	101	Guanajuato	place	pl
1	1	@lewellenmichael	Mexico	MX	mexico	2018-06-03	2018-05-31	19	-99	Mexico City		
2	2	@lewellenmichael	Mexico	MX	mexico	2017-11-05	2017-11-01	21	-86	Cancun		
3	3	@lewellenmichael	Jordan	JO	jordan	2017-08-07	2017-07-24	31	35	Amman		
4	4	@waylandchin	China	CN	china	2017-03-18	2017-02-17	40	122	Yingkou		

Preprocessing 1C

c. Filter the extreme observations, make sure minimum duration is 1 day!

In []:

```
#this give info about the quantiles
tripdata.describe()
```

Out[]:

	Unnamed: 0	latitude	longitude	duration
count	46510.000000	46510.000000	46510.000000	46282
mean	23254.500000	28.931477	11.912750	50 days 13:04:52.312346052
std	13426.424847	23.551976	78.767316	658 days 14:15:01.267546536
min	0.000000	-79.000000	-175.000000	-16831 days +00:00:00
25%	11627.250000	14.000000	-68.000000	3 days 00:00:00
50%	23254.500000	37.000000	9.000000	7 days 00:00:00
75%	34881.750000	47.000000	98.000000	23 days 00:00:00
max	46509.000000	78.000000	178.000000	74633 days 00:00:00

In []:

```
#code for 1c
# this will create a np array and give quantile in duration duration
extrem = pd.DataFrame(tripdata)
percentiles = np.array([0.05, 0.25, 0.50, 0.75, 0.975])
extrem.duration.quantile(percentiles)
#tripdata.groupby('duration').quantile(.68)
#my_DataFrame.groupby(['AGGREGATE']).agg({'duration': [q50, q90, 'max']})
```

Out[]:

```
0.050      1 days
0.250      3 days
0.500      7 days
0.750     23 days
0.975    261 days
Name: duration, dtype: timedelta64[ns]
```

In []:

```
# deleting the trips lower than 1 day and more than upper quantile
tripdata = tripdata[tripdata['duration'].dt.days >=1]
tripdata = tripdata[tripdata['duration'].dt.days <=660]
```

In []:

```
tripdata.duration.max()
```

Out []:

```
Timedelta('660 days 00:00:00')
```

In []:

```
tripdata['date_end'].dt.year.max()
```

Out []:

```
2106
```

In []:

```
tripdata = tripdata[tripdata['date_start'].dt.year > 2010]
```

In []:

```
tripdata.date_start.min()
```

Out []:

```
Timestamp('2011-01-01 00:00:00')
```

In []:

```
tripdata = tripdata[tripdata['date_end'].dt.year < 2022]
```

In []:

```
tripdata.date_end.max()
```

Out []:

```
Timestamp('2021-11-30 00:00:00')
```

In []:

```
tripdata.describe()
```

Out []:

	Unnamed: 0	latitude	longitude	duration
count	42309.000000	42309.000000	42309.000000	42309
mean	23080.457090	28.483183	13.540571	23 days 20:34:03.139757498
std	13470.733227	23.603071	78.985590	53 days 06:02:51.296091224
min	0.000000	-79.000000	-175.000000	1 days 00:00:00
25%	11389.000000	13.000000	-61.000000	3 days 00:00:00
50%	22988.000000	37.000000	11.000000	7 days 00:00:00
75%	34695.000000	47.000000	98.000000	23 days 00:00:00
max	46509.000000	78.000000	178.000000	660 days 00:00:00

Preprocessing 1D

d. Join the countrylist data to the trips data. Use the countrycode as a key.

In []:

```
#code for 1d
#tripdata.head()
#tripdata.country_code.value_counts()
#tripdata.country_code.unique()
#countrydata.alpha_2.value_counts()
```

```
trips_merged = pd.merge(tripdata, countrydata, left_on="country_code", right_on="alpha_2")
```

```
")
```

```
In [ ]:
```

```
countrydata.alpha_2.unique()
```

```
Out[ ]:
```

```
array(['AF', 'AX', 'AL', 'DZ', 'AS', 'AD', 'AO', 'AI', 'AQ', 'AG', 'AR',  
      'AM', 'AW', 'AU', 'AT', 'AZ', 'BS', 'BH', 'BD', 'BB', 'BY', 'BE',  
      'BZ', 'BJ', 'BM', 'BT', 'BO', 'BQ', 'BA', 'BW', 'BV', 'BR', 'IO',  
      'BN', 'BG', 'BF', 'BI', 'CV', 'KH', 'CM', 'CA', 'KY', 'CF', 'TD',  
      'CL', 'CN', 'CX', 'CC', 'CO', 'KM', 'CG', 'CD', 'CK', 'CR', 'CI',  
      'HR', 'CU', 'CW', 'CY', 'CZ', 'DK', 'DJ', 'DM', 'DO', 'EC', 'EG',  
      'SV', 'GQ', 'ER', 'EE', 'SZ', 'ET', 'FK', 'FO', 'FJ', 'FI', 'FR',  
      'GF', 'PF', 'TF', 'GA', 'GM', 'GE', 'DE', 'GH', 'GI', 'GR', 'GL',  
      'GD', 'GP', 'GU', 'GT', 'GG', 'GN', 'GW', 'GY', 'HT', 'HM', 'VA',  
      'HN', 'HK', 'HU', 'IS', 'IN', 'ID', 'IR', 'IQ', 'IE', 'IM', 'IL',  
      'IT', 'JM', 'JP', 'JE', 'JO', 'KZ', 'KE', 'KI', 'KP', 'KR', 'KW',  
      'KG', 'LA', 'LV', 'LB', 'LS', 'LR', 'LY', 'LI', 'LT', 'LU', 'MO',  
      'MG', 'MW', 'MY', 'MV', 'ML', 'MT', 'MH', 'MQ', 'MR', 'MU', 'YT',  
      'MX', 'FM', 'MD', 'MC', 'MN', 'ME', 'MS', 'MA', 'MZ', 'MM', nan,  
      'NR', 'NP', 'NL', 'NC', 'NZ', 'NI', 'NE', 'NG', 'NU', 'NF', 'MK',  
      'MP', 'NO', 'OM', 'PK', 'PW', 'PS', 'PA', 'PG', 'PY', 'PE', 'PH',  
      'PN', 'PL', 'PT', 'PR', 'QA', 'RE', 'RO', 'RU', 'RW', 'BL', 'SH',  
      'KN', 'LC', 'MF', 'PM', 'VC', 'WS', 'SM', 'ST', 'SA', 'SN', 'RS',  
      'SC', 'SL', 'SG', 'SX', 'SK', 'SI', 'SB', 'SO', 'ZA', 'GS', 'SS',  
      'ES', 'LK', 'SD', 'SR', 'SJ', 'SE', 'CH', 'SY', 'TW', 'TJ', 'TZ',  
      'TH', 'TL', 'TG', 'TK', 'TO', 'TT', 'TN', 'TR', 'TM', 'TC', 'TV',  
      'UG', 'UA', 'AE', 'GB', 'US', 'UM', 'UY', 'UZ', 'VU', 'VE', 'VN',  
      'VG', 'VI', 'WF', 'EH', 'YE', 'ZM', 'ZW'], dtype=object)
```

Part 2 People

People 2a

a. find how many people have at least high school diploma.

```
In [ ]:
```

```
#Count the different values in education raw, this sort on type
print(peopledata.education_raw.value_counts())
#summerize the string in that contains high in education
print("\npeople with High School diploma :")
#looking at what the string contains, and try to find High
print(peopledata.education_raw.str.contains("High").sum())

#Not using the string function you only get high school with one value and not multiple.
#len(peopledata.loc[peopledata['education_raw'] == 'High School'])
```

```
Bachelor's Degree          197
Master's Degree            115
High School                 58
High School, Bachelor's Degree  43
High School, Bachelor's Degree, Master's Degree  29
Bachelor's Degree, Master's Degree    9
Name: education_raw, dtype: int64
```

```
people with High School diploma :
130
```

People 2b

b. How many “Startun Founders” have attained a “Master’s Degree”?

In []:

```
# making a new data frame with name master_start and str contain with 2 condiftion
master_start = peopledata[peopledata.education_raw.str.contains("Master\s") & peopledat
a.work_raw.str.contains("Startup Founder")]

#take the length of this dataframe
len(master_start)
#Master_startup = peopledata[(peopledata.work_raw == 'Startup Founder') & (peopledata.edu
cation_raw == 'Masters Degree')]
```

Out[]:

53

People 2c

c. Who is the person with a Master’s Degree that has the highest number of followers? Bonus: Explore the individual further, what else can you find out?

In []:

```
#code for 2c
#her just sorting the values from max to low in followers
peopledata.sort_values(by=['followers'],ascending=False)

# the 2 person was having a master' degree
```

Out[]:

Unnamed: 0	username	followers	following	work_raw	education_raw
3096	@nomadlist_	2228	1	NaN	NaN
2043	@levelsio	2182	353	Software Dev, Startup Founder, Creative	High School, Bachelor's Degree, Master's Degree
3631	@vanschneider	520	13	NaN	NaN
494	@joelgascoigne	503	123	Startup Founder	NaN
1909	@marc	429	57	Web Dev, Software Dev, UI/UX Design, Startup F...	High School, Bachelor's Degree
...
2979	@cargix1	0	1	Software Dev	NaN
1620	@kutniy	0	2	NaN	NaN
2974	@mastinkipp	0	0	NaN	NaN
2971	@pahoda	0	0	NaN	NaN
3198	@tanjas	0	0	NaN	NaN

4016 rows x 6 columns

In []:

```
print(peopledata.iloc[2043])
```

```
Unnamed: 0          2043
username          @levelsio
followers          2182
following           353
work_raw          Software Dev, Startup Founder, Creative
education_raw    High School, Bachelor's Degree, Master's Degree
Name: 2043, dtype: object
```

Part 3 Trips

Trips 3a

a. Which country received the highest number of trips? – And which the lowest?

In []:

```
# code for 3a
#tripdata.head()
#tripdata.info()
#tripdata.shape

#save the value count in a new frame
max_travel = tripdata.country.value_counts()
print(max_travel)
```

```
United States      6771
Thailand           3255
United Kingdom     1977
Spain             1853
Germany           1800
...
45230              1
630 72             1
Mayotte            1
Andaman Sea        1
Lake Titicaca      1
Name: country, Length: 234, dtype: int64
```

the most travel to country is USA and lowest is some small islands

Trips 3b

b. Which region received the highest number of trips in 2017? Use the start of trips as a time reference.

In []:

```
tripdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 42309 entries, 0 to 46509
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Unnamed: 0            42309 non-null  int64   
 1   username              42309 non-null  object  
 2   country               42276 non-null  object  
 3   country_code          41796 non-null  object  
 4   country_slug          42276 non-null  object  
 5   date_end              42309 non-null  datetime64[ns]
 6   date_start            42309 non-null  datetime64[ns]
 7   latitude              42309 non-null  int64   
 8   longitude             42309 non-null  int64   
 9   place                 42309 non-null  object  
10   place_slug            42276 non-null  object  
11   duration              42309 non-null  timedelta64[ns]
dtypes: datetime64[ns](2), int64(3), object(6), timedelta64[ns](1)
memory usage: 4.2+ MB
```

In []:

```
trip_2017 = trips_merged.groupby('region')['date_start'].count()
```

```
#newframe.sort_values(['region', ['date_start'] <='2017-01-01'], ascending=[True, False],
inplace=True)
#newframe.loc[newframe["date_start"] <= "2017-01-01", "region"]
#newframe.sort_values(['region', 'date_start'], ascending=[True, False], inplace=True)

#newframe = newframe.groupby('region')['date_start'] <= '2017-01-01' #+ #df.groupby('away
_team')['away_score'].sum()
```

In []:

```
trips_2017 =trips_merged[trips_merged.date_end.dt.year<=2017]
trips_2017 =trips_2017[trips_2017.date_start.dt.year>=2017]
```

In []:

```
trips_2017.date_start.min()
```

Out[]:

```
Timestamp('2017-01-01 00:00:00')
```

In []:

```
trips_2017.date_end.max()
```

Out[]:

```
Timestamp('2017-12-31 00:00:00')
```

In []:

```
trips_2017 = trips_2017.groupby('region')['duration'].count()
```

In []:

```
trips_2017.sort_values(ascending=False)
```

Out[]:

```
region
Europe      4099
Asia        3255
Americas    3114
Africa       415
Oceania     357
Name: duration, dtype: int64
```

Most trips in Europe in 2017

Trips 3c

c. Which country in “Western Europe” did travelers spent least time?

In []:

```
#new frame is the merged frame with country and trips
trip_euro = trips_merged.groupby('sub_region')['duration'].count()
trip_euro.sort_values(ascending=False)
trip_eurowest = trips_merged[trips_merged.sub_region == 'Western Europe']
trip_eurowest.sort_values(by="duration", ascending=False)
trip_eurowest.groupby("country")["duration"].sum()
```

Out[]:

```
country
Austria      5331 days
Belgium      4473 days
France       40115 days
Germany      49991 days
Liechtenstein    1 days
```



```
Liechtenstein      1 days
Luxembourg         3350 days
Monaco              64 days
Netherlands        35199 days
Switzerland        14283 days
Name: duration, dtype: timedelta64[ns]
```

people spend the least time in liechtenstein on ly 1 day

```
In [ ]:
```

```
trip_eurowest = trip_eurowest.groupby("country")["duration"].sum()
```

```
In [ ]:
```

```
trip_eurowest.plot(kind='bar', title="Western Europe days in total")
```

Trips 3d

d. Do nomad Startup Founders tend to have shorter or longer trips on average?

```
In [ ]:
```

```
people_merged = peopledata.merge(trips_merged, how='left', on="username")
```

```
#peoplemerge = peopledata(newframe, how='left', on="username")
```

```
#not_startup_founders = peopledata[~peopledata['work_raw'].str.contains('Startup Founder')]
```

```
In [ ]:
```

```
people_merged.duration.mean()
```

```
Out[ ]:
```

```
Timedelta('24 days 22:20:12.908085206')
```

```
In [ ]:
```

```
people_SF = people_merged[people_merged['work_raw'].str.contains('Startup Founder', na=False)]
```

```
In [ ]:
```

```
people_SF.duration.mean()
```

```
Out[ ]:
```

```
Timedelta('23 days 13:02:42.688296639')
```

Start up Founders tend to have shorter trips

Trips 3e

e. visualize over-time median trip duration overall (bonus: and split by world-region) The plot will look weird ^^.
PyHint: Resample by week ('W') and calculate the size of observations. **RHint:** Use the floor_date function to reset dates by week

```
In [ ]:
```

```
import seaborn as sns
import altair as alt
```

```
sns.set(color_codes = True)
```

```
import matplotlib.pyplot as plt
```

```
In [ ]:
```

```
people_merged['duration'] = people_merged['duration'].dt.days
```

```
In [ ]:
```

```
people_merged= people_merged[people_merged.date_start.dt.year<=2018]
```

```
In [ ]:
```

```
people_mg =people_merged.groupby(by=[people_merged.date_start.dt.year, people_merged.region]) ['duration'].median()
```

```
In [ ]:
```

```
print(people_mg)
```

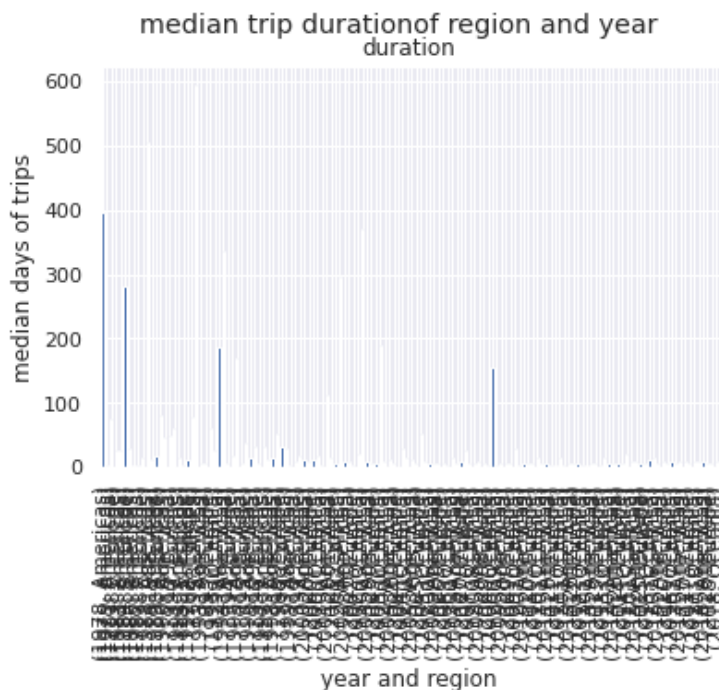
```
date_start  region
1978        Americas    395.0
1979        Americas     3.0
1981        Americas    76.5
          Europe     27.0
1983        Americas    27.0
          ...
2018        Africa     9.0
          Americas     8.0
          Asia        7.0
          Europe      6.0
          Oceania     9.0
Name: duration, Length: 138, dtype: float64
```

```
In [ ]:
```

```
people_mg.plot(stacked=True, kind='bar', subplots=True, ylabel='median days of trips', title='median trip durationof region and year', xlabel='year and region')
```

```
Out[ ]:
```

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f087eb48b90>],
      dtype=object)
```



filter issue, it included all of the years

```
In [ ]:
```

```
people_mgr = people_mgr.reset_index()
```

In []:

```
people_mgr.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 138 entries, 0 to 137
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date_start  138 non-null   int64  
 1   region      138 non-null   object  
 2   duration    138 non-null   float64
dtypes: float64(1), int64(1), object(1)
memory usage: 3.4+ KB
```

In []:

```
sns.lineplot(data=people_mgr, x='date_start', y='duration', hue='region')
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f087ea454d0>



In [2]:

```
jupyter nbconvert --to html Assignment.ipynb
```

```
File "<ipython-input-2-89c24475e90b>", line 1
  jupyter nbconvert --to html Assignment.ipynb
    ^
```

SyntaxError: invalid syntax